

Executive Memo: Early-Tenure Retention Trigger

Recommendation to pilot a proactive retention intervention for the highest-churn customer segment

The Ask

Approve a 4-6 week pilot of a proactive retention trigger targeting 966 month-to-month customers in their first 6 months with no support add-on — the single highest-churn segment identified (65.7% churn vs. 26.5% company-wide).

Why Now

- This segment alone represents \$63.3K/month (\$760K annualized) in monthly-recurring revenue at elevated risk.
- Churn concentrates early: 53.3% of customers in months 0-6 churn, vs. 9.5% by months 49-72 — the intervention window is well-defined and short.
- Customers with a TechSupport add-on churn at 15.2% vs. 41.6% without — a 63.6% relative gap that this feature is designed to test and close.

Projected Impact (Pilot-Scale, First-Pass Estimate)

Metric	Value
Target segment size	966 customers
Assumed adoption rate	40%
Customers retained (modeled)	~161
Annualized revenue saved	~\$127K
Intervention cost	~\$3.1K
Modeled ROI	~40x

These figures come from a first-pass sizing model with two explicit assumptions (adoption rate, cost per outreach) — see the impact model and PRD Risks section. The pilot's purpose is to replace those assumptions with real numbers before any full rollout decision.

What We're Asking For

- Sign-off to build the trigger logic and CS tooling flag (Phase 1, ~2 weeks).
- CS bandwidth for a randomized 50/50 treatment-vs-control pilot (Phase 2, ~4-6 weeks).
- A follow-up decision point after the pilot to approve or reject full rollout, based on measured — not modeled — churn reduction.

Key Risk

The core assumption — that giving a support add-on reduces churn, rather than merely correlating with customers who were already going to stay — is unproven. The pilot's control group exists specifically to test this before we commit further budget.